

10

SOLUTIONS. Solutions vary; this was a great example from a student.

Dr. Ian Frommer

Data Science Techniques 11 February 2026

### Data Science Techniques – Homework 4

#### Question #1

- Example 1: `csvcut -c City,Country world_famous_places_2024.csv | csvlook`

City	Country
Daxis	France
New York City	United States
Paris	France
Beijing/Multiple	China
Agra	India
Rome	Italy
New York City	United States
Sydney	Australia
Cusco Region	Peru
Beijing	China
Siem Reap	Cambodia
Paris	France
New York City	United States
Las Vegas	United States
New York City	United States
San Francisco	United States
Washington D.C.	United States
Barcelona	Spain
Anaheim	United States
Orlando	United States
Arizona	United States
Cairo	Egypt
Pisa	Italy
Versailles	France
Athens	Greece
Dubai	United Arab Emirates
London	United Kingdom
London	United Kingdom
London	United Kingdom
Rio de Janeiro	Brazil

In Example 1, I select the columns “City” and “Country” from the `world_famous_places_2024.csv` using `csvcut`, which filters .csv files. From there, I piped it to the `csvlook` command, which displays the .csv file in tabular form. This complete command allows us to view specific portions of data in an organized fashion.

- Example 2: `csvstat -freq world_famous_places_2024.csv | sort`

```

1. Place_Name: { "Eiffel Tower": 1, "Times Square": 1, "Louvre Museum": 1, "Great Wall of China": 1, "Taj Mahal": 1 }
2. Country: { "United States": 10, "France": 8, "United Kingdom": 3, "China": 2, "Italy": 2 }
3. City: { "New York City": 4, "Paris": 3, "London": 3, "Beijing/Multiple": 1, "Agra": 1 }
4. Annual_Visitors_Millions: { "15": 2, "6": 2, "7": 1, "50": 1, "8.7": 1 }
5. Type: { "Historic Monument": 3, "Monument/Statue": 2, "Archaeological Site": 2, "Historic Palace": 2, "Cathedral": 2 }
6. UNESCO_World_Heritage: { "True": 17, "False": 13 }
7. Year_Built: { "1951": 2, "1889": 1, "1904": 1, "1793": 2, "226 BC - 1644 AD": 1 }
8. Entry_Fee_USD: { "0": 9, "35": 2, "20": 2, "22": 1, "10": 1 }
9. Best_Visit_Month: { "Apr-May/Sept-Oct": 5, "Apr-June/Sept-Oct": 5, "May-Sept": 4, "Apr-June/Sept-Nov": 2, "Oct-March": 2 }
10. Region: { "North America": 10, "Western Europe": 7, "Southern Europe": 4, "East Asia": 2, "South America": 2 }
11. Tourism_Revenue_Million_USD: { "180": 3, "70": 2, "120": 2, "65": 2, "35": 2 }
12. Average_Visit_Duration_Hours: { "1.5": 5, "1": 5, "2": 4, "3": 4, "4": 3 }
13. Famous_For: { "Iconic iron lattice tower, symbol of Paris": 1, "Bright lights, Broadway shows, New Year's Eve ball drop": 1, "World's most visited museum, home to Mona Lisa": 1, "Ancient defensive structure visible from space": 1, "White marble mausoleum, symbol of love": 1 }

```

Nice  
The command used here, `csvstat`, provides a descripted summary of columns, but by adding “`-freq`” afterwards, the most common values for each column are outputted. I piped them to the `sort` command to order them neatly.

- Example 3: `csvcut -c Place_Name,Year_Built,Famous_For world_famous_places_2024.csv | csvgrep -c Year_Built -r '^\\d{4}$' | csvsort -c Year_Built | csvlook`

Place_Name	Year_Built	Famous_For
Angkor Wat	1,150	Largest religious monument, Hindu-Buddhist temple
Notre-Dame Cathedral	1,345	Gothic masterpiece, medieval Catholic cathedral
Leaning Tower of Pisa	1,372	Tilted bell tower, architectural anomaly
Forbidden City	1,420	Imperial palace of Ming and Qing dynasties
Machu Picchu	1,450	Ancient Incan citadel in the Andes mountains
Taj Mahal	1,653	White marble mausoleum, symbol of love
Palace of Versailles	1,682	Opuent royal residence, Hall of Mirrors
Buckingham Palace	2,703	Official residence of British monarch
Louvre Museum	1,793	World's most visited museum, home to Mona Lisa
Central Park	1,857	Urban park oasis in Manhattan
Big Ben	1,859	Clock tower, Westminster Palace
Statue of Liberty	1,886	Symbol of freedom and democracy
Eiffel Tower	1,889	Iconic iron lattice tower, symbol of Paris
Tower Bridge	1,894	Victorian Gothic suspension bridge
Times Square	1,904	Bright lights, Broadway shows, New Year's Eve ball drop
Las Vegas Strip	1,905	Casino resorts, entertainment, nightlife
Lincoln Memorial	1,922	Memorial to President Abraham Lincoln
Empire State Building	1,931	Art Deco skyscraper, NYC icon
Christ the Redeemer	1,931	Art Deco statue of Jesus Christ
Golden Gate Bridge	1,937	Suspension bridge, engineering marvel
Disneyland (California)	1,958	Original Disney theme park
Magic Kingdom (Orlando)	1,971	Disney World's flagship park with Cinderella Castle
Sydney Opera House	1,973	Unique sail-like design, performing arts center
Burj Khalifa	2,818	Tallest building in the world
Great Pyramid of Giza	2,550	Only remaining ancient wonder, pharaoh's tomb

As displayed in the image above, I used `csvcut` to splice the .csv file and extract the three columns listed. From there, I used the `csvgrep` command to filter out any years that did not match the four-digit format as specified in the regular expression. Then, I sorted all entries by "Year\_Built" from oldest to newest. `csvlook` outputs this data in a more readable way.

- Example 4: `csvgrep -c Region -m 'North America' world_famous_places_2024.csv | csvcut -C Region,Tourism_Revenue_Million_USD | csvsort -r -c Annual_Visitors_Millions | csvlook`

Place_Name	Country	City	Annual_Visitors_Millions	Type	UNESCO_World_Heritage
Times Square	United States	New York City	50.0	Urban Landmark	False
Central Park	United States	New York City	42.0	Park	False
Las Vegas Strip	United States	Las Vegas	41.7	Entertainment District	False
Magic Kingdom (Orlando)	United States	Orlando	17.0	Theme Park	False
Disneyland (California)	United States	Anaheim	16.0	Theme Park	False
Golden Gate Bridge	United States	San Francisco	15.0	Bridge	False
Lincoln Memorial	United States	Washington D.C.	8.5	Monument/Memorial	False
Grand Canyon	United States	Arizona	6.0	Natural Wonder	True
Empire State Building	United States	New York City	4.6	Skyscraper	False
Statue of Liberty	United States	New York City	4.3	Monument/Statue	True

Year_Built	Entry_Fee_USD	Best_Visit_Month	Average_Visit_Duration_Hours	Famous_For
1904	8	Apr-June/Sept-Nov	1.5	Bright lights, Broadway shows, New Year's Eve ball drop
1957	6	May-Sept	1.0	Urban park oasis in Manhattan
1965	6	March-May/Sept-Nov	3.0	Casino resorts, entertainment, nightlife
1971	100	Jan-May/Sept-Dec	10.0	Disney World's flagship park with Cinderella Castle
1955	100	Apr-May/Sept-Nov	8.0	Original Disney theme park
1937	8	Sept-Nov	1.0	Suspension bridge, engineering marvel
1922	8	Apr-June/Sept-Oct	1.0	Memorial to President Abraham Lincoln
Natural Formation	35	March-May/Sept-Nov	5.0	Massive canyon carved by Colorado River
1931	44	May-Sept	1.5	Art Deco skyscraper, NYC icon
1986	25	May-Sept	2.0	Symbol of freedom and democracy

In Example 4, I filtered for all sites in the North American region using `csvgrep` before cutting the two columns listed from the final output. Then, I sorted these remaining entries by "Annual\_Visitors\_Millions" to order them from most to least popular. Finally, I ran the `csvlook` command to improve the output appearance.

## Question #2

- A. I used the feline.db for this portion of the homework. I used the following query to gather the first and last names, as well as the total number of orders, of each female customer.

```
sqlite> SELECT "First Name", "Last Name", COUNT(Orders."Order ID") AS "Orders"
...> FROM Customers INNER JOIN Orders ON Customers."Customer ID" = Orders."Customer ID"
...> WHERE Customers.Gender = 'F'
...> GROUP BY Customers."Last Name";
Victoria|Bailey|2
Bernadette|Chapman|1
Hannah|Mitchell|2
Abigail|Russell|2
Sonia|Wright|2
sqlite>
```

- B. In addition to .tables, I used the .schema command to identify the columns in the Customer and Orders tables.

```
sqlite> .schema Customers
CREATE TABLE IF NOT EXISTS "Customers" (
    "Customer ID" INTEGER,
    "First Name" TEXT,
    "Last Name" TEXT,
    "Email Address" TEXT,
    "Street Address" TEXT,
    "City" TEXT,
    "State" TEXT,
    "Zip" INTEGER,
    "Gender" TEXT,
    "Credit Card Number" INTEGER
);
```

- I used the .quit command to exit out of sqlite3 whenever I needed to either switch databases or make other changes in the command line.

```
sqlite> .quit
(CsvTools) PS C:\Users\Spring 26\dst_course_work\Homeworks\Homework_4> sqlite
```

To change the output appearance, I used the .mode list and .headers on commands prior to running the query shown to get the output below.

```
sqlite> .mode list
sqlite> .headers on
sqlite> SELECT * FROM Orders;
Order ID|Customer ID|Date Ordered|Date Delivered
470|501|1/5/2018|1/15/2018
471|502|1/5/2018|1/15/2018
472|503|1/5/2018|1/15/2018
473|504|1/5/2018|1/15/2018
474|505|1/5/2018|1/15/2018
475|506|1/5/2018|1/15/2018
476|507|1/5/2018|1/15/2018
477|508|1/5/2018|1/15/2018
478|509|1/5/2018|1/15/2018
480|501|2/1/2018|4/1/2018
481|502|3/15/2018|4/3/2018
482|503|2/26/2018|4/19/2018
483|504|3/18/2018|4/3/2018
484|505|3/31/2018|4/28/2018
485|506|2/10/2018|3/31/2018
486|507|2/3/2018|4/3/2018
487|508|3/2/2018|4/10/2018
488|509|1/31/2018|3/30/2018
489|510|3/1/2018|4/16/2018
sqlite>
```